

Rithujaa Rajendrakumar

+1 201-204-6526 | rithujaa128@gmail.com | [LinkedIn](#) | [Github](#) | [Portfolio](#)

TECHNICAL SKILLS

Programming & Querying: Python, SQL, PySpark, R

Data Engineering & Cloud: AWS (S3, Glue, Lambda, EC2), Azure (Data Factory, Databricks, Synapse), Snowflake, Apache Airflow, Apache Spark, Kafka, ETL/ELT, Data Warehousing, Orchestration

Machine Learning & Statistics: Regression, Classification, Clustering, NLP, Feature Engineering, XGBoost, Random Forest, Statistical Analysis

AI & Automation: LLMs (OpenAI/Anthropic APIs), embeddings, prompt engineering, AI-assisted workflow automation

Tools & BI: Tableau, Power BI, Streamlit, Git, Docker, Microsoft Excel

EXPERIENCE

Data Science Intern

Sept 2025 – Dec 2025

Trilemma Foundation

Remote

- Built a 6-agent ensemble strategy for dynamic Bitcoin accumulation, achieving a 90.37% win rate over uniform DCA across 3,075 rolling 12-month backtesting windows
- Engineered features from 369 variables spanning on-chain metrics, macroeconomic indicators, and NLP-based sentiment signals (FinBERT, RoBERTa), applying statistical analysis to drive allocation decisions
- Built an interactive Streamlit dashboard to visualize agent performance, sentiment trends, and rolling-window backtest results

Analytical Engineer

May 2023 – Aug 2024

Sonline LLC

Remote

- Developed and optimized SQL workflows to automate reporting pipelines across client datasets, cutting query turnaround time from days to hours
- Built interactive Tableau dashboards and data visualizations for 3+ enterprise clients, surfacing pipeline health and key business metrics
- Worked directly with the founding team and stakeholders across 2–3 enterprise accounts to translate business requirements into analytical solutions, presenting findings in client-facing meetings
- Cleaned and validated large client datasets using SQL and Excel, resolving recurring data inconsistencies that improved reporting accuracy

PROJECTS

AI-Powered Data Integration Pipeline | *Python, LLMs, Sentence Transformers, HDBSCAN, DuckDB*

- Built an AI system to unify 1M+ records across 3 cities with incompatible schemas, eliminating weeks of manual mapping
- Used LLMs, sentence embeddings, and HDBSCAN clustering to automatically group and match equivalent data points across datasets
- Evaluated output quality using precision/recall metrics to ensure reliable results before downstream use

Airbnb Occupancy & Revenue Prediction | *Python, scikit-learn, XGBoost, CatBoost, VADER*

- Built predictive ML models (XGBoost, Random Forest, CatBoost) across 11K+ NYC Airbnb listings, achieving $R^2 = 0.65$ for revenue and $AUC = 0.94$ for top-earner classification
- Engineered 83 features (VADER sentiment, KNN imputation, outlier treatment); translated findings into actionable pricing and operations recommendations

Real-Time Data Streaming Pipeline | *AWS EC2, Apache NiFi, Amazon S3, Snowflake, Snowpipe, Streams, Tasks*

- Streamed generated event data from EC2 (Python/Faker) through NiFi into S3, modeling data into Raw / Current / History layers with Snowpipe, Streams, and Tasks for automated CDC

Apache Kafka Real-Time Pipeline | *Apache Kafka, Python, AWS S3, Docker*

- Built a real-time streaming pipeline with Kafka to ingest and process trip data through custom producers and consumers, persisting records to S3 and containerizing the stack with Docker Compose

EDUCATION

New York University, Centre for Data Science

New York, NY

Master of Science in Data Science (100% CDS Pathbreaker Scholarship), GPA 3.6/4

Sep 2024 – May 2026

University of British Columbia

Kelowna, Canada

Bachelor of Science in Data Science, GPA 3.7/4

Sep 2019 – May 2023